

In [181]:

```
'''
Author: Zhuofan Dong
Date: 10/13/2022
Subject: Assignment 1 for Statistical Analysis
'''
```

Out[181]: '\n\nAuthor: Zhuofan Dong\nDate: 10/13/2022\nSubject: Assignment 1 for Statistical Analysis\n\n'

In [21]:

```
import numpy as np
import pandas as pd
import os
os.getcwd()
os.chdir(r'C:\Users\Primo\OneDrive\Desktop')
data = pd.read_csv('cars.csv')
```

In [22]:

```
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 428 entries, 0 to 427
Data columns (total 15 columns):
 #   Column              Non-Null Count  Dtype  
---  --
 0   Make                428 non-null   object 
 1   Model               428 non-null   object 
 2   Type                428 non-null   object 
 3   Origin              428 non-null   object 
 4   DriveTrain          428 non-null   object 
 5   MSRP                428 non-null   int64  
 6   Invoice              428 non-null   int64  
 7   EngineSize          428 non-null   float64 
 8   Cylinders            426 non-null   float64 
 9   Horsepower          428 non-null   int64  
10  MPG_City             428 non-null   int64  
11  MPG_Highway          428 non-null   int64  
12  Weight               428 non-null   int64  
13  Wheelbase            428 non-null   int64  
14  Length              428 non-null   int64  
dtypes: float64(2), int64(8), object(5)
memory usage: 50.3+ KB
```

Question 1

In [33]:

```
data['MSRP'].describe().apply("{0:.2f}".format)
```

Out[33]:

```
count      428.00
mean       32774.86
std        19431.72
min        10280.00
25%        20334.25
50%        27635.00
75%        39205.00
max        192465.00
Name: MSRP, dtype: object
```

```
In [34]: data['Weight'].describe().apply("{0:.2f}".format)
```


```
Out[34]: count      428.00  
mean      3577.95  
std       758.98  
min      1850.00  
25%      3104.00  
50%      3474.50  
75%      3977.75  
max      7190.00  
Name: Weight, dtype: object
```

```
In [35]: data['Length'].describe().apply("{0:.2f}".format)
```

```
Out[35]: count      428.00  
mean      186.36  
std       14.36  
min      143.00  
25%      178.00  
50%      187.00  
75%      194.00  
max      238.00  
Name: Length, dtype: object
```

Question 2

```
In [67]: import matplotlib.pyplot as plt  
  
df_cylinder = data['Cylinders']  
df_cylinder = df_cylinder.dropna()  
  
freq_cylinder = df_cylinder.value_counts()  
  
freq_values = freq_cylinder.tolist()  
freq_index = freq_cylinder.index.tolist()  
freq_index = [int(x) for x in freq_index]  
  
plt.bar(freq_index, freq_values)  
plt.xticks(freq_index, freq_index)  
plt.show()
```



Yes, it looks like the binomial distribution.

Question 3

```
In [102... df_msrp = data['MSRP']
df_weight = data['Weight']
df_length = data['Length']

## Formulas from Module

def univariate (y):
    # Initialize
    y_nvalid = 0
    y_min = None
    y_max = None
    y_mean = None
    # Loop through all the elements
    for u in y:
        if (not np.isnan(u)):
            y_nvalid = y_nvalid + 1
            if (y_min is not None):
                if (u < y_min):
                    y_min = u
            else:
                y_min = u
            if (y_max is not None):
                if (u > y_max):
                    y_max = u
            else:
                y_max = u
            if (y_mean is not None):
                y_mean = y_mean + u
            else:
                y_mean = u
    # Finalize
    if (y_nvalid > 0):
        y_mean = y_mean / y_nvalid
    return (y_nvalid, y_min, y_max, y_mean)

def shimazaki_criterion (y, d_list):
    number_bins = []
    matrix_boundary = []
    shimazaki_criterion = []
    y_nvalid, y_min, y_max, y_mean = univariate (y)
    if (y_nvalid <= 0):
        raise ValueError('There are no non-missing values in the data vector.')
    else:
        # Loop through the bin width candidates
        for delta in d_list:
            y_middle = delta * np.round(y_mean / delta)
            n_bin_left = np.ceil((y_middle - y_min) / delta)
            n_bin_right = np.ceil((y_max - y_middle) / delta)
            y_low = y_middle - n_bin_left * delta
            # Assign observations to bins starting from 0
            list_boundary = []
            n_bin = n_bin_left + n_bin_right
            bin_index = 0
            bin_boundary = y_low
```

```

list_boundary.append(bin_boundary)
for i in np.arange(n_bin):
    bin_boundary = bin_boundary + delta
    bin_index = np.where(y > bin_boundary, i+1, bin_index)
    list_boundary.append(bin_boundary)
    # Count the number of observations in each bins
    uvalue, ucount = np.unique(bin_index, return_counts = True)
    # Calculate the average frequency
    mean_ucount = np.mean(ucount)
    ssd_ucount = np.mean(np.power((ucount - mean_ucount), 2))
    criterion = (2.0 * mean_ucount - ssd_ucount) / delta / delta
    number_bins.append(n_bin)
    matrix_boundary.append(list_boundary)
    shimazaki_criterion.append(criterion)

return(number_bins, matrix_boundary, shimazaki_criterion)

```

```

In [127]: #Next bin width is 200000 which is greater than the range
d_list_msrp = [1, 2, 2.5, 5, 10, 20, 25, 50, 100, 200, 250, 500, 1000, 2000, 2500, 5000]
#MSRP Binwidth
msrp_number_bins, msrp_matrix_boundary, msrp_shimazaki_criterion = shimazaki_criterion(d_list_msrp)

df_msrp_binwidth = pd.DataFrame(list(zip(msrp_number_bins, msrp_shimazaki_criterion)),

#Return the index of min value
minValueIndex = df_msrp_binwidth['Shamazaki Criterion'].idxmin()
minValueIndex

```

Out[127]: 5000.0

```

In [126]: #Next bin width is 100 which is greater than the range
d_list_length = [1, 2, 2.5, 5, 10, 20, 25, 50]
#Length Binwidth
length_number_bins, length_matrix_boundary, length_shimazaki_criterion = shimazaki_criterion(d_list_length)

df_length_binwidth = pd.DataFrame(list(zip(length_number_bins, length_shimazaki_criterion)),

#Return the index of min value
minValueIndex = df_length_binwidth['Shamazaki Criterion'].idxmin()
minValueIndex

```

Out[126]: 2.5

```

In [125]: #Next bin width is 10000 which is greater than the range
d_list_weight = [1, 2, 2.5, 5, 10, 20, 25, 50, 100, 200, 250, 500, 1000, 2000, 2500, 5000]
#Weight Binwidth
weight_number_bins, weight_matrix_boundary, weight_shimazaki_criterion = shimazaki_criterion(d_list_weight)

df_weight_binwidth = pd.DataFrame(list(zip(weight_number_bins, weight_shimazaki_criterion)),

#Return the index of min value
minValueIndex = df_weight_binwidth['Shamazaki Criterion'].idxmin()
minValueIndex


```

Out[125]: 100.0

According to the dataframes respectively, we can find the minimum shamazaki criterion corresponding to the bin width. For MSRP, the smallest criterion is when bin width is 5000. For Length, the smallest criterion is when bin width is 2.5. For Weight, the smallest criterion is when bin width is 100.

Histogram for MSRP


```
In [138... for delta, bin_boundary in zip(d_list_msrp, msrp_matrix_boundary):
    if delta == 5000:
        plt.figure(figsize = (4,3), dpi = 200)
        plt.hist(df_msrp, bins = bin_boundary, align = 'mid')
        plt.title('Binwidth = ' + str(delta))
        plt.xlabel('MSRP')
        plt.ylabel('Number of Observations')
        plt.grid(axis = 'y')
        plt.show()
    else:
        continue
```



Histogram for Length

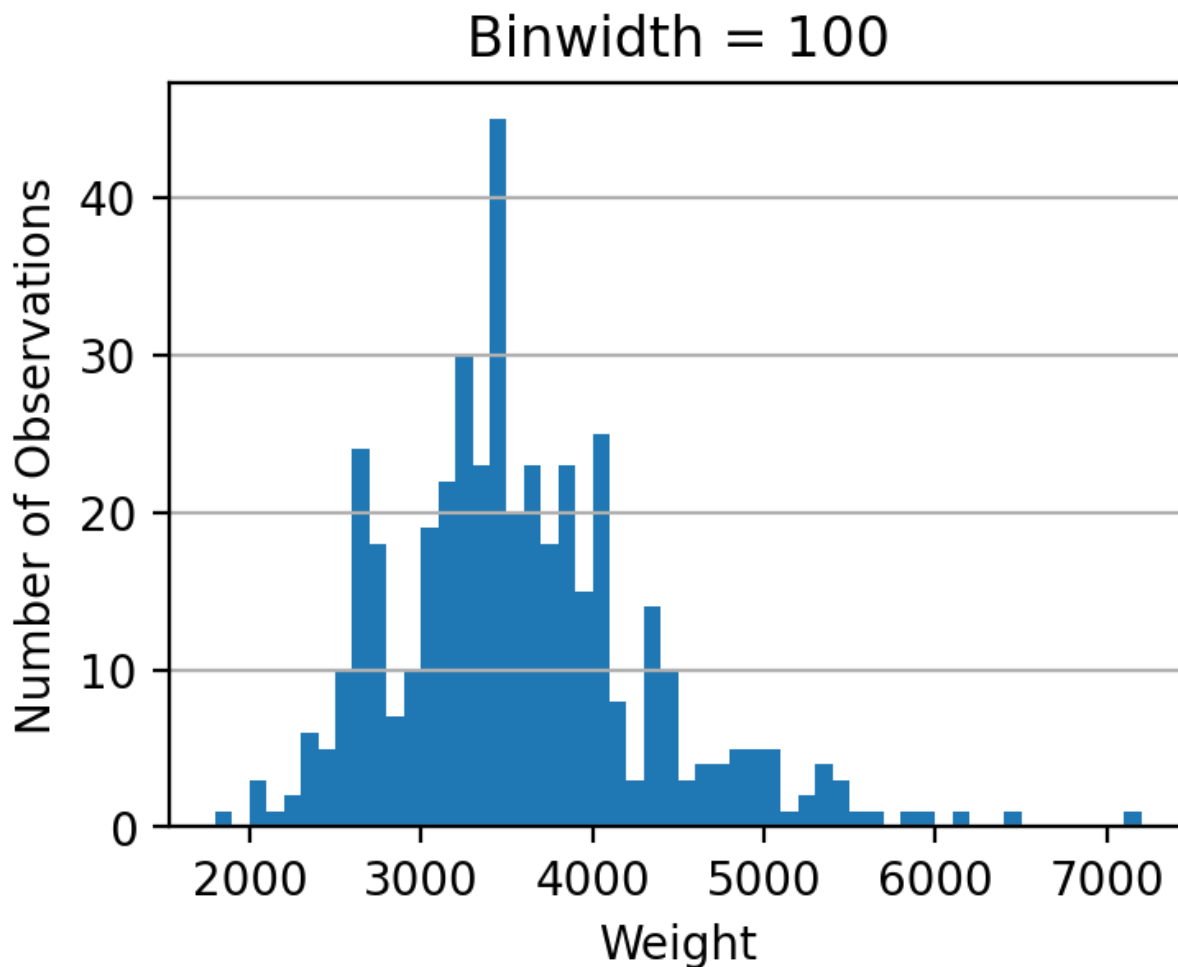
```
In [131... for delta, bin_boundary in zip(d_list_length, length_matrix_boundary):
    if delta == 2.5:
        plt.figure(figsize = (4,3), dpi = 200)
        plt.hist(df_length, bins = bin_boundary, align = 'mid')
        plt.title('Binwidth = ' + str(delta))
```

```
plt.xlabel('Length')
plt.ylabel('Number of Observations')
plt.grid(axis = 'y')
plt.show()
else:
    continue
```



Histogram for Weight

```
In [137... for delta, bin_boundary in zip(d_list_weight, weight_matrix_boundary):
    if delta == 100:
        plt.figure(figsize = (4,3), dpi = 200)
        plt.hist(df_weight, bins = bin_boundary, align = 'mid')
        plt.title('Binwidth = ' + str(delta))
        plt.xlabel('Weight')
        plt.ylabel('Number of Observations')
        plt.grid(axis = 'y')
        plt.show()
    else:
        continue
```



Question 4

In [139... `from scipy.stats import norm, shapiro, anderson`


Analysis for Length

```
In [145... n = len(df_length)
# Sort the values in ascending order
df_length_sorted = np.sort(df_length)
# Calculate the hypothetical quantiles
u_mean = np.mean(df_length)
u_stddev = np.std(df_length)
p = np.arange(1,(n+1)) / (n + 0.5)
z = norm.ppf(p, loc = u_mean, scale = u_stddev)
print('\nZ Minimum = ', np.min(z))
print('Z Maximum = ', np.max(z))
fig, ax = plt.subplots(nrows = 1, ncols = 1, dpi = 200, figsize = (8,8))
ax.scatter(df_length_sorted, z, marker = 'o', c = 'royalblue')
ax.set_xlim(xmin = 140, xmax = 240)
ax.set_ylim(ymin = 140, ymax = 240)
ax.set_xticks(np.arange(140, 240, 5.0))
ax.set_yticks(np.arange(140, 240, 5.0))
ax.axline((200,200), slope = 1.0, linestyle = '--', color = 'red')
ax.set_xlabel('Observation Quantile')
```

```
ax.set_ylabel('Hypothetical Quantile')
ax.set_aspect(1.0)
ax.set_title('Normal Distribution')
ax.margins(x = 0.1, y = 0.1)
ax.grid()
plt.show()
```

Z Minimum = 145.78897900509048

Z Maximum = 230.01831948486614



```
In [143... # Shapiro-Wilks test for normality
shapiro_test = shapiro(df_length)
print('\nShapiro Test = ', shapiro_test[0])
print('      p-value = ', shapiro_test[1])
# Anderson-Darling test for normality
anderson_test = anderson(df_length, dist = 'norm')
print('\n Anderson Test = ', anderson_test[0])
print('Critical Values = ', anderson_test[1])
print('      p-values = ', anderson_test[2]/100.0)
```


Shapiro Test = 0.9911824464797974
 p-value = 0.011863326653838158

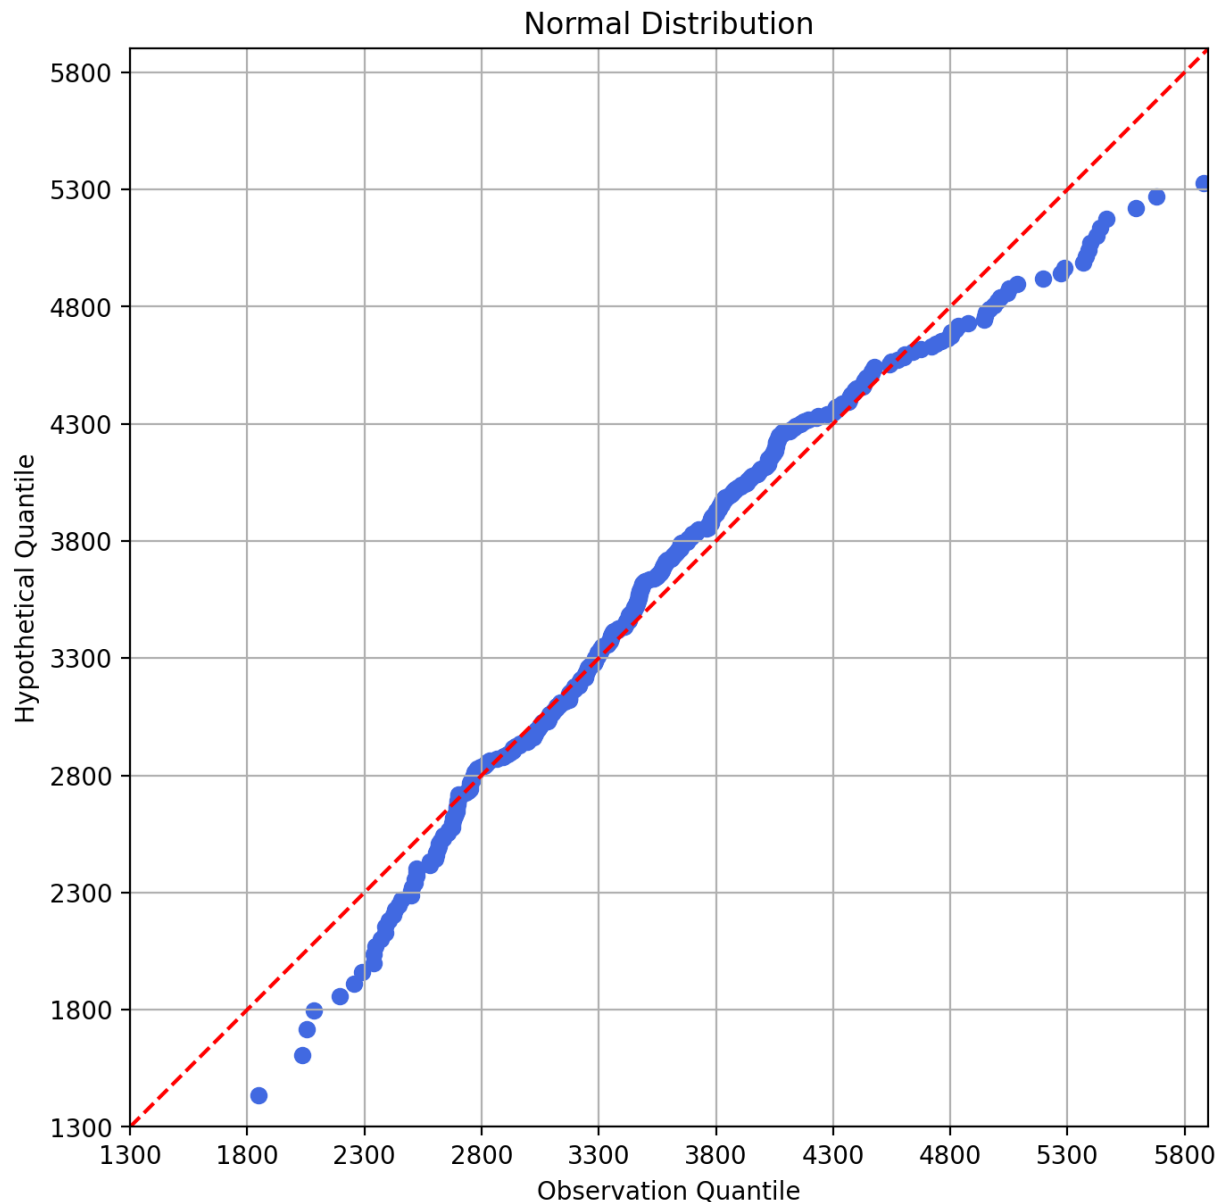
Anderson Test = 1.286219586060156
 Critical Values = [0.571 0.65 0.78 0.91 1.082]
 p-values = [0.15 0.1 0.05 0.025 0.01]

It is obvious that both Shapiro and Anderson Test indicate that we have to reject the null hypothesis, given the p-value is less than 0.05 significance level for Shapiro Test and test statistic is greater than the critical value of any level for Anderson Test

Analysis for Weight

```
In [151... n = len(df_weight)
# Sort the values in ascending order
df_weight_sorted = np.sort(df_weight)
# Calculate the hypothetical quantiles
u_mean = np.mean(df_weight)
u_stddev = np.std(df_weight)
p = np.arange(1,(n+1)) / (n + 0.5)
z = norm.ppf(p, loc = u_mean, scale = u_stddev)
print('\nZ Minimum = ', np.min(z))
print('Z Maximum = ', np.max(z))
fig, ax = plt.subplots(nrows = 1, ncols = 1, dpi = 200, figsize = (8,8))
ax.scatter(df_weight_sorted, z, marker = 'o', c = 'royalblue')
ax.set_xlim(xmin = 1400, xmax = 5900)
ax.set_ylim(ymin = 1400, ymax = 5900)
ax.set_xticks(np.arange(1300, 6000, 500.0))
ax.set_yticks(np.arange(1300, 6000, 500.0))
ax.axline((200,200), slope = 1.0, linestyle = '--', color = 'red')
ax.set_xlabel('Observation Quantile')
ax.set_ylabel('Hypothetical Quantile')
ax.set_aspect(1.0)
ax.set_title('Normal Distribution')
ax.margins(x = 0.1, y = 0.1)
ax.grid()
plt.show()
```

Z Minimum = 1433.1995347295897
 Z Maximum = 5885.678607741062



```
In [152... # Shapiro-Wilks test for normality
shapiro_test = shapiro(df_weight)
print('\nShapiro Test = ', shapiro_test[0])
print('      p-value = ', shapiro_test[1])
# Anderson-Darling test for normality
anderson_test = anderson(df_weight, dist = 'norm')
print('\n Anderson Test = ', anderson_test[0])
print('Critical Values = ', anderson_test[1])
print('      p-values = ', anderson_test[2]/100.0)
```

```
Shapiro Test = 0.9589154720306396
p-value = 1.4718520846557226e-09
```

```
Anderson Test = 3.804003476182743
Critical Values = [0.571 0.65 0.78 0.91 1.082]
p-values = [0.15 0.1 0.05 0.025 0.01 ]
```

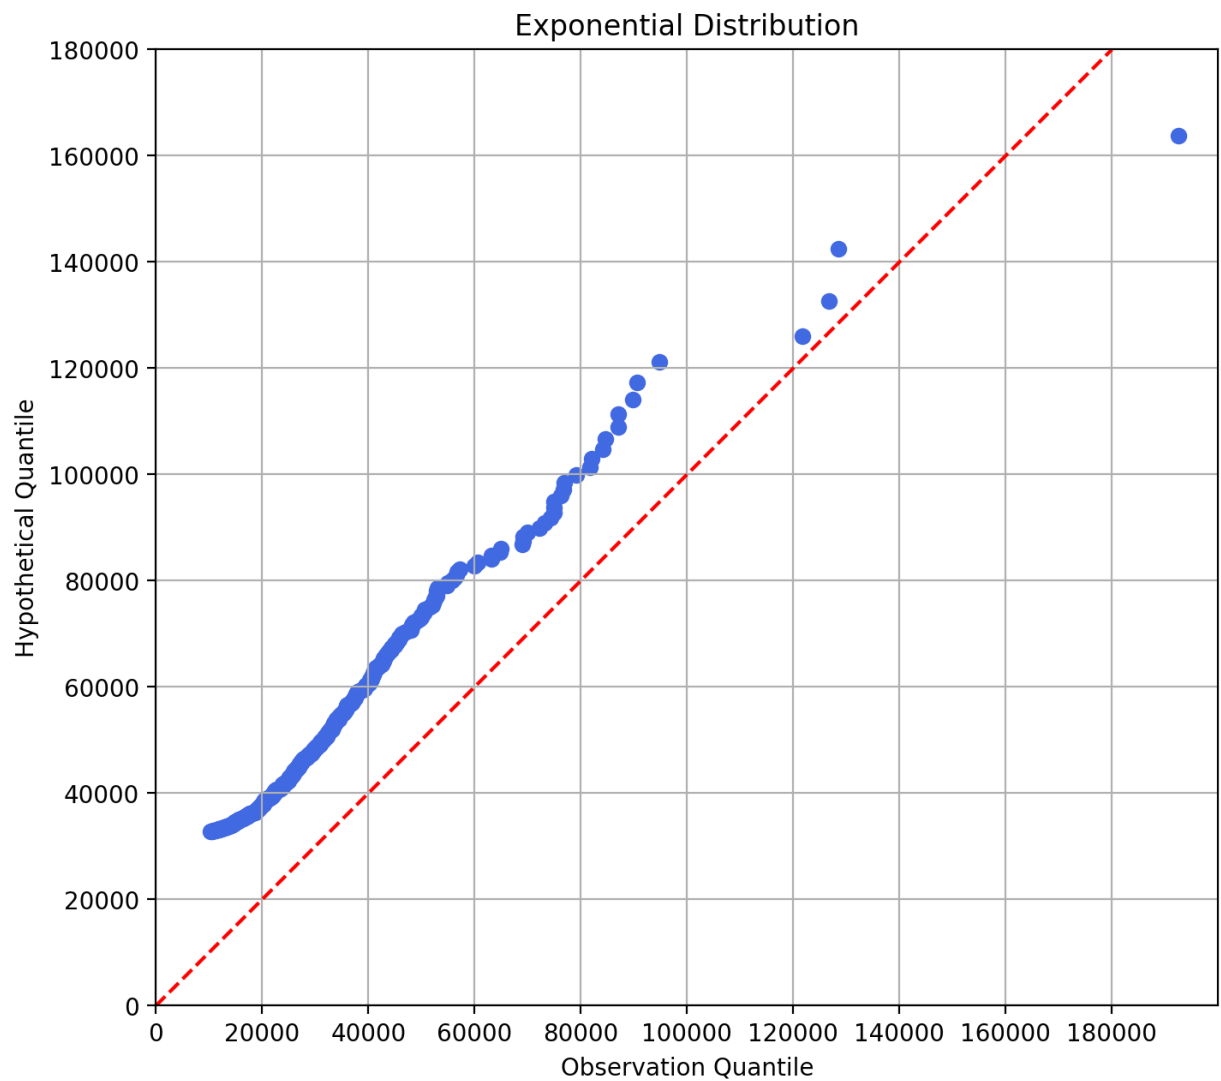
It is obvious that both Shapiro and Anderson Test indicate that we have to reject the null hypothesis, given the p-value is less than 0.05 significance level for Shapiro Test and test statistic is greater than the critical value of any level for Anderson Test

Question 5

In [180... `from scipy.stats import expon`

```
n = len(df_msrp)
# Sort the values in ascending order
df_msrp_sorted = np.sort(df_msrp)
# Calculate the hypothetical quantiles
u_mean = np.mean(df_msrp)
u_stddev = np.std(df_msrp)
p = np.arange(1, (n+1)) / (n + 0.5)
z = expon.ppf(p, loc = u_mean, scale = u_stddev)
print('\nZ Minimum = ', np.min(z))
print('Z Maximum = ', np.max(z))
fig, ax = plt.subplots(nrows = 1, ncols = 1, dpi = 200, figsize = (8,8))
ax.scatter(df_msrp_sorted, z, marker = 'o', c = 'royalblue')
ax.set_xlim(xmin = 0, xmax = 200000)
ax.set_ylim(ymin = 0, ymax = 20000)
ax.set_xticks(np.arange(0, 200000, 20000.0))
ax.set_yticks(np.arange(0, 20000, 2000.0))
ax.axline((200,200), slope = 1.0, linestyle = '--', color = 'red')
ax.set_xlabel('Observation Quantile')
ax.set_ylabel('Hypothetical Quantile')
ax.set_aspect(1.0)
ax.set_title('Exponential Distribution')
ax.margins(x = 0.1, y = 0.1)
ax.grid()
plt.show()
```

```
Z Minimum = 32820.20329818783
Z Maximum = 163852.35057753872
```



```
In [179... print("The mean of MSRP is ", u_mean, " and the standard deviation of MSRP is ", u_std)
```

The mean of MSRP is 32774.85514018692 and the standard deviation of MSRP is 19409.002794915406

Given the Q-Q plot and the comparison of mean and the standard deviation of the MSRP data, I do not think it follows the exponential distribution

```
In [ ]:
```